

Bayesian Data Analysis - Assignment 4

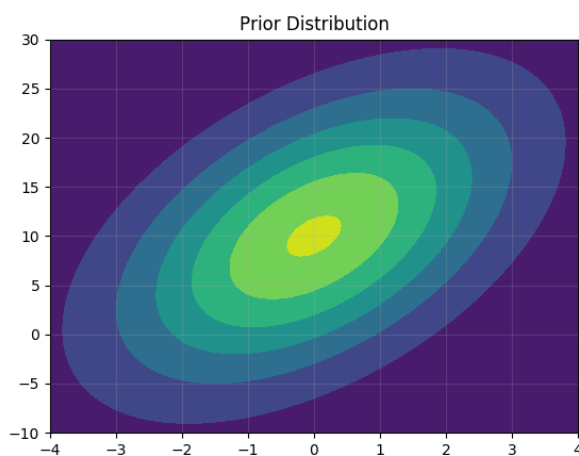
October 7, 2018

We can derive σ , μ and *covariance matrix* from the given normal distributions of $\alpha \sim N(0, 2^2)$ and $\beta \sim N(10, 10^2)$.

```
1 sigma_a = 2
2 sigma_b = 10
3 mu_a = 0
4 mu_b = 10
5 cor = 0.5
6 cov_matrix = np.array([
7     [sigma_a**2,          cor * sigma_a * sigma_b],
8     [cor * sigma_a * sigma_b,  sigma_b**2]]
9 )
```

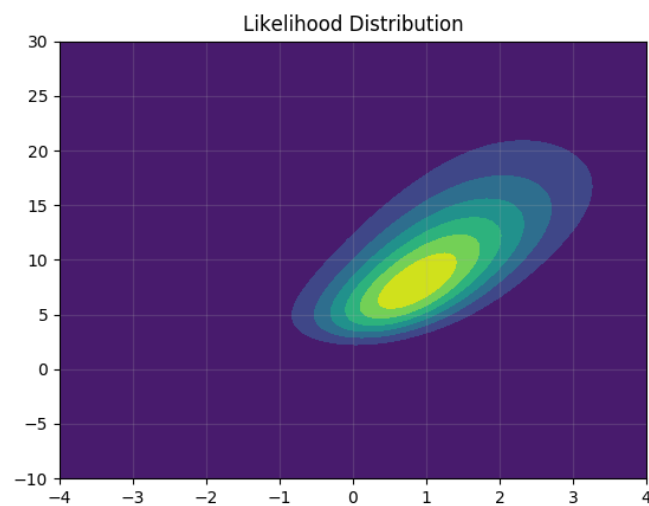
This initialization allows us to calculate and plot **prior distribution**:

```
1 # create a grid and it's points using x and y ranges
2 alpha, beta = np.meshgrid(np.linspace(-4, 4, 100), np.linspace(-10, 30, 100))
3 points = np.dstack((alpha, beta))
4
5 # prior distribution
6 mean = np.array([mu_a, mu_b])
7 prior_multivar_nor = stats.multivariate_normal(mean, cov_matrix)
8 prior = prior_multivar_nor.pdf(points)
9 plt.contourf(alpha, beta, prior)
```



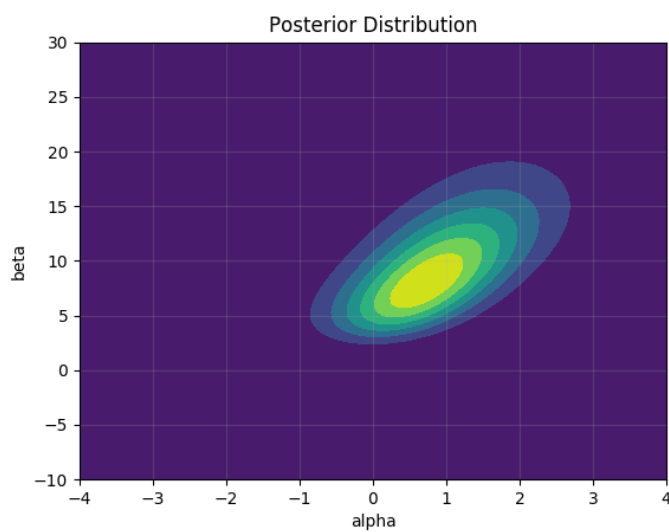
To calculate the **likelihood** though, we have to take into consideration the data from the book that indicates the values of *doses*, *deaths* and *number of animals*; and use the `bioassaylp` function:

```
1 doses = np.array([-0.86, -0.3, -0.05, 0.72])
2 deaths = np.array([0, 1, 3, 5])
3 number_of_animals = np.array([5, 5, 5, 5])
4
5 likelihood = bioassaylp(alpha, beta, doses, deaths, number_of_animals)
6 plt.contourf(alpha, beta, np.exp(likelihood))
```



Calculating **posterior** is straight forward, we have to just multiply the likelihood with prior:

```
1 posterior = prior * np.exp(likelihood)
2 plt.contourf(alpha, beta, posterior)
```



Using the prior distribution we can **draw a sample with a size = 10000**:

```
1 samples_from_prior = prior_multivar_nor.rvs(10000)
```

We can then compute **the importance ratios for each draw** like:

```
1 logit = 1 / (
2     1 + np.exp(
3         -(samples_from_prior[:, 0, None] + samples_from_prior[:, 1, None] * doses)
4     )
5 )
6
7 weights_of_likelihood = np.prod(
8     logit**deaths * (1 - logit)**(number_of_animals - deaths),
9     axis=1
10 )
11 print('Shape of the weights of the likelihood: ', weights_of_likelihood.shape)
```

Consequently, **the effective sample size** S_{eff} would be calculated as:

```
1 weights_norm = (weights_of_likelihood) / np.sum(weights_of_likelihood)
2 S_eff = 1 / np.sum(weights_norm**2)
3 print('The effective sample size: ', round(S_eff, 4))
```

Which outputs:

```
1 $ The effective sample size: 2725.6818
```

We can then easily compute the **posterior mean** using importance sampling:

```
1 mean_posterior = sum(
2     weights_of_likelihood[ : , None] * samples_from_prior
3 ) / sum(weights_of_likelihood)
4 print('The posterior mean of alpha : ', round(mean_posterior[0], 4))
5 print('The posterior mean of beta : ', round(mean_posterior[1], 4))
```

Which are:

```
1 $ The posterior mean of alpha : 0.9786
2 $ The posterior mean of beta : 10.4321
```

Using importance resampling we can obtain a **posterior sample of size 1000** of alpha and beta:

```
1 gen_nums = np.random.RandomState(0).choice(
2     a=10000, size=1000, replace=False, p=weights_norm
3 )
4 resamples = samples_from_prior[gen_nums]
5 print('The posterior mean of alpha: ', round(np.mean(resamples[:, 0]), 4))
6 print('The posterior mean of beta: ', round(np.mean(resamples[:, 1]), 4))
```

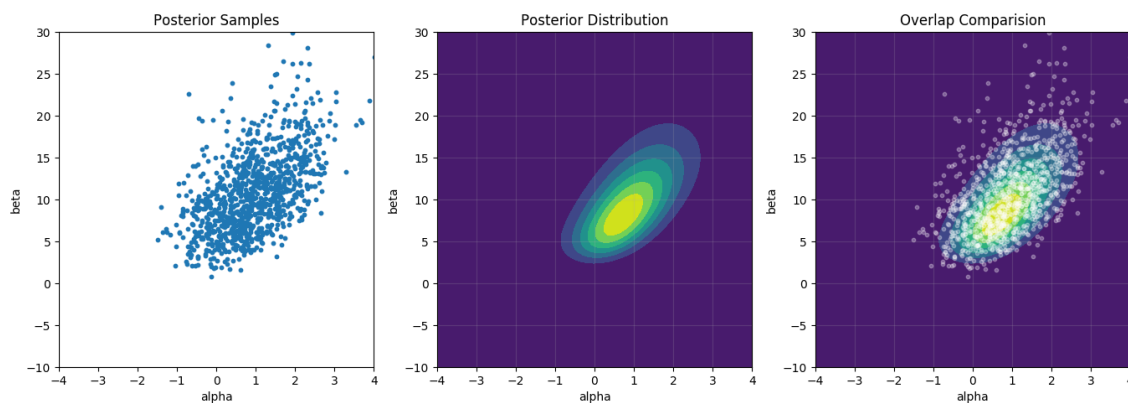
Which are:

```
1 $ The posterior mean of alpha: 0.9359
2 $ The posterior mean of beta: 10.3614
```

If we combine these two we will have $\alpha = [0.9359, 0.9786]$ and $\beta = [10.4321, 10.3614]$.

Finally, we can **plot a scatterplot of the obtained posterior sample** from importance resampling and compare it to the heatmap:

```
1 fig, axes = plt.subplots(1, 3, figsize=(16, 5))
2 axes[0].scatter(resamples[:, 0], resamples[:, 1], 10)
3 axes[1].grid(linewidth=0.9, alpha=0.2)
4 axes[2].contourf(alpha, beta, posterior)
5 axes[2].scatter(resamples[:, 0], resamples[:, 1], 10, alpha=0.3, color='w')
```



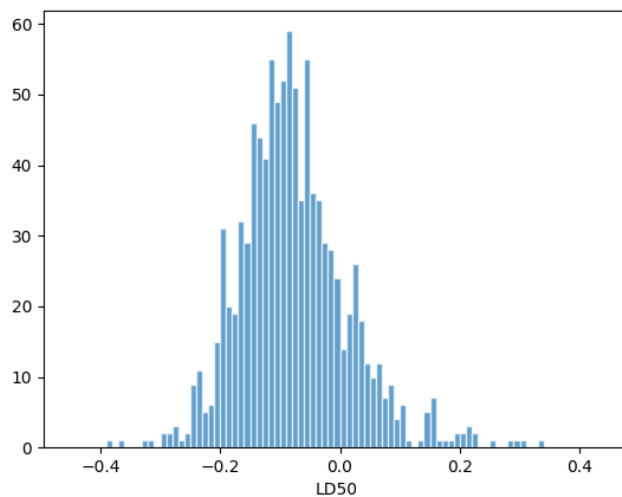
The estimate $p(\beta > 0 | x, n, y)$ probability can be calculated as:

```
1 beta_resample = resamples[:, 1]
2 alpha_resample = resamples[:, 0]
3 pos = beta_resample > 0
4 probab_drug_harmful = (beta_resample[pos].size / (beta_resample.size))
5 print('Probability of the drug being harmful: ', round(probab_drug_harmful, 3))
```

```
1 $ Probability of the drug being harmful is very close to: 1
```

And the **histogram** can be plotted like:

```
1 sample_ld50 = -alpha_resample[pos] / beta_resample[pos]
2 y_range = np.arange(-0.45, 0.45, 0.01)
3 plt.hist(sample_ld50, y_range, ec='white', alpha=0.7)
```



Appendix A Source code

```

1  import matplotlib
2  matplotlib.use('TkAgg')
3  import matplotlib.pyplot as plt
4  from mpl_toolkits.mplot3d import Axes3D
5  from scipy import stats
6  import numpy as np
7
8  def bioassaylp(a, b, x, y, n):
9      # last axis for the data points
10     a = np.expand_dims(a, axis=-1)
11     b = np.expand_dims(b, axis=-1)
12     # these help using chain rule in derivation
13     t = a + b*x
14     et = np.exp(t)
15     z = et/(1.+et)
16     # negative log posterior (error function to be minimized)
17     lp = np.sum(y*np.log(z)+ (n-y)*np.log(1.0-z), axis=-1)
18     return lp
19
20 # Init all the params based on the description
21 sigma_a = 2
22 sigma_b = 10
23 mu_a = 0
24 mu_b = 10
25 cor = 0.5
26 cov_matrix = np.array([
27     [sigma_a**2,          cor * sigma_a * sigma_b],
28     [cor * sigma_a * sigma_b,  sigma_b**2]]
29 )
30
31 # create a grid and it's points using x and y ranges
32 alpha, beta = np.meshgrid(np.linspace(-4, 4, 100), np.linspace(-10, 30, 100))

```

```

33 points = np.dstack((alpha, beta))
34
35 # prior distribution
36 mean = np.array([mu_a, mu_b])
37 prior_multivar_nor = stats.multivariate_normal(mean, cov_matrix)
38 prior = prior_multivar_nor.pdf(points)
39
40 plt.contourf(alpha, beta, prior)
41 plt.title('Prior Distribution')
42 plt.grid(linewidth=0.9, alpha=0.2)
43 plt.savefig('./ex4/report/1_prior_distribution.png')
44 plt.figure()
45
46 # likelihood
47 doses = np.array([-0.86, -0.3, -0.05, 0.72])
48 deaths = np.array([0, 1, 3, 5])
49 number_of_animals = np.array([5, 5, 5, 5])
50
51 likelihood = bioassaylp(alpha, beta, doses, deaths, number_of_animals)
52 plt.contourf(alpha, beta, np.exp(likelihood))
53 plt.title('Likelihood Distribution')
54 plt.grid(linewidth=0.9, alpha=0.2)
55 plt.savefig('./ex4/report/2_likelihood_distribution.png')
56 plt.figure()
57
58 '''
59 1. Calculate the posterior density in a grid of points around the prior
60 (alpha: 0 +- 4, beta: 10 +- 20) and plot a heatmap of the density.
61 '''
62 posterior = prior * np.exp(likelihood)
63 plt.contourf(alpha, beta, posterior)
64 plt.xlabel('alpha')
65 plt.ylabel('beta')
66 plt.title('Posterior Distribution')
67 plt.grid(linewidth=0.9, alpha=0.2)
68 plt.savefig('./ex4/report/3_posterior_distribution.png')
69 plt.figure()
70
71 '''
72 2. Sample draws of alpha and beta from the prior distribution.
73 '''
74 samples_from_prior = prior_multivar_nor.rvs(10000)
75 print('Shape of the samples from prior: ', samples_from_prior.shape)
76
77 '''
78 3. Compute the importance ratios for each draw (target distribution is the posterior).
79 '''
80 logit = 1 / (
81     1 + np.exp(
82         -(samples_from_prior[:, 0, None] + samples_from_prior[:, 1, None] * doses)
83     )
84 )
85
86 weights_of_likelihood = np.prod(
87     logit**deaths * (1 - logit)**(number_of_animals - deaths),

```

```

88     axis=1
89 )
90 print('Shape of the weights of the likelihood: ', weights_of_likelihood.shape)
91
92 '''
93 4. Using the importance ratios, compute the effective sample size  $S_{eff}$  and report it.
94 If  $S_{eff}$  is less than 1000, repeat the importance sampling with more draws.
95 '''
96 weights_norm = (weights_of_likelihood) / np.sum(weights_of_likelihood)
97 S_eff = 1 / np.sum(weights_norm**2)
98 print('The effective sample size: ', round(S_eff, 4))
99
100 '''
101 5. Compute the posterior mean using importance sampling and report it.
102 '''
103 mean_posterior = sum(
104     weights_of_likelihood[ : , None] * samples_from_prior
105 ) / sum(weights_of_likelihood)
106 print('The posterior mean of alpha : ', round(mean_posterior[0], 4))
107 print('The posterior mean of beta : ', round(mean_posterior[1], 4))
108
109 '''
110 6. Use importance resampling to obtain a posterior sample of size 1000
111 of alpha and beta.
112 '''
113 gen_nums = np.random.RandomState(0).choice(
114     a=10000, size=1000, replace=False, p=weights_norm
115 )
116 resamples = samples_from_prior[gen_nums]
117 print('Mean alpha: ', round(np.mean(resamples[:, 0]), 4))
118 print('Mean beta: ', round(np.mean(resamples[:, 1]), 4))
119
120 '''
121 7. Using the posterior sample obtained via importance resampling:
122 7.1 Plot a scatterplot of the obtained posterior sample and compare that to
123 the heatmap plotted earlier.
124 '''
125 fig, axes = plt.subplots(1, 3, figsize=(16, 5))
126 axes[0].set_xlim([-4, 4])
127 axes[0].set_ylim([-10, 30])
128 axes[0].set_xlabel('alpha')
129 axes[0].set_ylabel('beta')
130 axes[0].set_title('Posterior Samples')
131 axes[0].scatter(resamples[:, 0], resamples[:, 1], 10)
132
133 axes[1].set_xlim([-4, 4])
134 axes[1].set_ylim([-10, 30])
135 axes[1].set_xlabel('alpha')
136 axes[1].set_ylabel('beta')
137 axes[1].set_title('Posterior Distribution')
138 axes[1].contourf(alpha, beta, posterior)
139 axes[1].grid(linewidth=0.9, alpha=0.2)
140
141 axes[2].set_xlim([-4, 4])
142 axes[2].set_ylim([-10, 30])

```

```

143 axes[2].set_xlabel('alpha')
144 axes[2].set_ylabel('beta')
145 axes[2].set_title('Overlap Comparision')
146 axes[2].contourf(alpha, beta, posterior)
147 axes[2].scatter(resamples[:, 0], resamples[:, 1], 10, alpha=0.3, color='w')
148 axes[2].grid(linewidth=0.9, alpha=0.2)
149
150 plt.subplots_adjust(left=0.1, right=0.9, top=0.9, bottom=0.1)
151 plt.savefig('./ex4/report/4_posterior_comparision.png')
152 plt.figure()
153
154 '''
155     7.2 Report an estimate for  $p(\beta > 0|x, n, y)$ , that is, the
156     probability that the drug is harmful.
157 '''
158 beta_resample = resamples[:, 1]
159 alpha_resample = resamples[:, 0]
160 pos = beta_resample > 0
161 probab_drug_harmful = (beta_resample[pos].size/(beta_resample.size))
162 print('Probability of the drug being harmful: ', round(probab_drug_harmful, 3))
163
164 '''
165     7.3 Draw a histogram of the draws from the posterior distribution of the LD50
166     conditional on  $\beta > 0$  (see Figure 3.4).
167 '''
168 sample_ld50 = -alpha_resample[pos]/beta_resample[pos]
169 y_range = np.arange(-0.45, 0.45, 0.01)
170 plt.hist(sample_ld50, y_range, ec='white', alpha=0.7)
171
172 plt.xlabel('LD50')
173 plt.savefig('./ex4/report/5_histogram.png')
174 plt.figure()

```