

Bayesian Data Analysis - Assignment 7

November 5, 2018

1 Question

Before analyzing the predictions, it will be beneficial to plot the data and look at the trend. We can do that simply by:

```
plt.plot(years, drowning)
z = np.polyfit(years, drowning, 1)
trend = np.poly1d(z)
plt.plot(years, trend(years), 'r--')
plt.savefig('./ex7/report/drowning.png')
```




Figure 1: Number of people drown per year in Finland.

From the trend we can see that the number of drown in Finland is decreasing.

1.1 Stan code

The *stan* code described in the question 1 description had 2 main issues:

- There was no lower bound for *sigma*. It was declared as *real sigma*. This is not good, as it might lead to some miscalculations. The correct declaration is:

```
real<lower=0> sigma
```

- The second issue was in *generated quantities*:

```
ypred=normal_rng(mu, sigma)
```

That is incorrect because we have to evaluate the prediction on years, which is *xpred*. The correct declaration is:

```
ypred = normal_rng(alpha + beta * xpred, sigma);
```

1.2 Numerical value of τ

- The numerical value of τ is:

$$\tau = 26.78888\dots$$

It was calculated by putting various values into *scale* and checking the statement *dist.cdf*(-69):

```
dist = norm(loc=0, scale=26.78)
print(dist.cdf(-69))
```

The value of *dist.cdf*(-69) should be 0.1/2 for the correct τ .

1.3 Prior in the stan model

- The initial value of both α and β in stan model were taken from uniform prior. That was changed for β to weekly-informative prior:

```
beta ~ normal(0, tau);
```

Please look at *Appendix A Source code for Question 1* for details.

1.4 Predictions for 2019

After we have built our stan model, we can predict any upcoming year simply by writing:

```
data = dict(
  N=len(years),
  x=years,
  y=drowning,
  xpred=2019,
  tau=26.78,
)
fit = stan_model.sampling(data=data)
y_pred = fit.extract()['ypred']
plt.hist(y_pred, bins=20, ec='white')
```

And that will give us the predictions as:

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
alpha	1809.3	23.27	827.88	177.23	1269.3	1797.3	2371.9	3436.3	1266	1.0
beta	-0.84	0.01	0.41	-1.65	-1.12	-0.83	-0.57	-0.02	1266	1.0
sigma	26.4	0.09	3.24	20.86	24.15	26.11	28.36	33.52	1376	1.0
mu[1]	152.33	0.22	8.56	135.69	146.53	152.26	158.08	168.97	1544	1.0
mu[2]	151.5	0.21	8.21	135.56	145.94	151.5	157.07	167.4	1576	1.0
mu[3]	150.66	0.2	7.86	135.32	145.31	150.66	155.96	165.95	1614	1.0
...
mu[36]	123.06	0.21	8.41	106.21	117.51	123.01	128.62	139.86	1597	1.0
ypred	121.19	0.54	28.65	64.71	101.84	120.69	140.59	177.37	2824	1.0

As we can see from the above table the the *ypred* is the prediction for 2019 which is 121 drowning.




Figure 2: Histogram of prediction for 2019.

2 Question

2.1 Pooled model

In the pooled model all the machines are considered as one entity, thus we have to combine the all the measurements into one and perform our prediction on the whole data; rather than a subset. With that in mind, we have to first read the data and then flatten it into one array:

```
machines = pd.read_fwf('./ex7/factory.txt', header=None).values
machines_pooled = machines.flatten()
```

The stan model for this is stated in the *Appendix B Source code for Question 2*. Next, we can use it to make the predictions:

```
fit_pooled = model_pooled.sampling(data=data_pooled)
```

Which gives us this table:

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
mu	92.87	0.06	3.35	86.22	90.65	92.87	95.06	99.41	2748	1.0
sigma	18.77	0.05	2.55	14.6	16.94	18.5	20.32	24.68	2468	1.0
ypred	93.33	0.32	19.08	55.19	80.93	93.76	105.58	130.87	3550	1.0
lp_	-99.3	0.02	0.99	-102.0	-99.66	-99.01	-98.6	-98.35	1759	1.0

The prediction histogram looks like this:




Figure 3: Prediction histogram with pooled model.

As we combined all the measurements into one, the μ value will be the same for sixth machine, seventh machine or all the machines combined. As mentioned before we don't treat each machine as a separate entity, that's exactly the reason why all the μ will be equal. This is how the histogram looks like:




Figure 4: μ histogram with pooled model for 6th, 7th or all machines combined.

2.2 Separate model

In the separate model we treat every machine as an individual entity, thus when calculating σ or μ we take into consideration only a single machine. The combination of all machines should

not effect σ or μ . The stan model for this is stated in the *Appendix B Source code for Question 2*

- *the posterior distribution of the mean of the quality measurements of the sixth machine:*
This can be calculated simply by:

```
fit_separate = model_separate.sampling(data=data_separate, n_jobs=-1)
print(fit_separate)
```

This gives us a nice table where we can allocate the μ of the sixth machine:

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
mu[1]	75.93	0.36	15.3	44.6	68.25	75.98	83.95	106.05	1833	1.0
mu[2]	106.57	0.34	10.33	88.75	101.66	106.3	110.89	126.97	944	1.0
mu[3]	87.8	0.23	9.49	67.93	83.03	87.92	92.65	106.74	1646	1.0
mu[4]	111.56	0.17	5.91	99.43	108.6	111.52	114.46	123.49	1225	1.0
mu[5]	89.91	0.27	9.27	70.6	85.84	90.0	94.32	107.29	1143	1.0
mu[6]	86.45	0.37	14.33	58.26	78.54	86.19	93.8	117.74	1477	1.0
sigma[1]	30.62	0.55	20.17	13.08	19.47	25.44	34.81	79.54	1352	1.0
sigma[2]	18.56	0.37	11.89	7.69	11.64	15.26	21.23	52.47	1017	1.0
sigma[3]	19.64	0.31	12.34	8.41	12.64	16.55	22.55	49.52	1574	1.0
sigma[4]	12.03	0.2	6.95	4.96	7.6	10.09	14.15	30.48	1265	1.0
sigma[5]	16.98	0.39	11.8	7.06	10.52	13.68	19.27	46.61	907	1.0
sigma[6]	30.16	0.5	18.0	12.28	18.87	25.17	35.34	79.14	1289	1.0
ypred	86.0	0.7	37.13	12.32	67.29	85.77	104.62	162.64	2827	1.0
lp__	-81.11	0.1	2.99	-88.04	-82.89	-80.75	-78.92	-76.34	913	1.0

Which is **86.45**. We can also plot the histogram of it:

```
mu_data_separate = fit_separate.extract()['mu']
plt.hist(mu_data_separate[:, 5], bins=20, ec='white')
```




Figure 5: μ histogram with separate model for 6th machine.

- *the predictive distribution for another quality measurement of the sixth machine:*
The predictive distribution can be extracted from *ypred*.

```
y_pred_separate = fit_separate.extract()['ypred']
plt.hist(y_pred_separate, bins=20, ec='white')
```




Figure 6: Prediction histogram with separate model for the sixth machine.

- *the posterior distribution of the mean of the quality measurements of the seventh machine:* As it was stated before, in the separate model we treat each machine separately. Consequently, we have no any information about the seventh machine. Thus we cannot tell anything about its posterior distribution.

2.3 Hierarchical model

The hierarchical model is quite interesting. It does treat every machine as a separate entity, but also computes the combination of all the machines as one entity. For that reason it can predict measurements for the machines which have no data. For example, there is no data about the seventh machine, but this model can predict its posterior distribution. The stan model for this is stated in the *Appendix B Source code for Question 2*.

- *the posterior distribution of the mean of the quality measurements of the sixth machine:* The same logic, as in separate model, follows here. We start simply by:

```
fit_hierarchical = model_hierarchical.sampling(data=data_hierarchical, n_jobs=-1)
print(fit_hierarchical)
```

Which again, gives us a nice table where we can allocate the μ of the sixth machine:

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
mu0	92.13	0.36	9.72	71.8	88.13	92.64	97.03	108.53	738	1.0
sigma0	17.08	0.46	12.29	5.38	10.49	14.32	20.03	45.45	729	1.01
mu[1]	79.71	0.15	6.56	66.76	75.32	79.65	84.13	92.72	1992	1.0
mu[2]	103.07	0.14	6.48	90.04	98.92	103.04	107.31	116.52	2247	1.0
mu[3]	88.89	0.09	6.18	76.41	84.87	88.85	92.92	101.25	4231	1.0
mu[4]	107.56	0.18	6.68	94.24	103.04	107.69	112.16	120.04	1345	1.0
mu[5]	90.47	0.1	6.12	78.28	86.59	90.45	94.46	102.45	3904	1.0
mu[6]	87.32	0.1	6.25	75.07	83.08	87.31	91.64	99.36	3647	1.0
sigma	15.18	0.06	2.27	11.55	13.55	14.91	16.54	20.41	1572	1.0

ypred6	87.0	0.26	16.89	52.64	76.11	86.87	98.52	120.0	4076	1.0
mu7	91.77	0.59	24.43	44.47	82.15	92.6	103.4	133.87	1687	1.0
lp__	-108.9	0.09	2.52	-114.8	-110.3	-108.4	-107.0	-105.3	755	1.0

Which is **87.32** and it is quite close to what we had with the separate model: **86.45**. We can now plot the histogram:

```
mu_data_hierarchical = fit_hierarchical.extract()['mu']
plt.hist(mu_data_hierarchical[:, 5], bins=20, ec='white')
```




Figure 7: μ histogram with hierarchical model for 6th machine.

- *the predictive distribution for another quality measurement of the sixth machine:*
The prediction can also be extracted from the above table by:

```
y_pred_hierarchical = fit_hierarchical.extract()['ypred6']
plt.hist(y_pred_hierarchical, bins=20, ec='white')
```




Figure 8: Prediction histogram with hierarchical model for 6th machine.

- *the posterior distribution of the mean of the quality measurements of the seventh machine:*
Referring to the table we printed above, we can say that the μ for the seventh machine is

91.77. We can now draw the posterior distribution of the mean for the seventh machine simply by:

```
mu_data_hierarchical_7 = fit_hierarchical.extract()['mu7']
plt.hist(mu_data_hierarchical_7, bins=20, ec='white')
```




Figure 9: μ histogram with hierarchical model for 7th machine.

Appendix A Source code for Question 1

```
1 import matplotlib
2 matplotlib.use('TkAgg')
3
4 from scipy.stats import norm
5 import matplotlib.pyplot as plt
6 import numpy as np
7 import pandas as pd
8 import pystan
9
10 drowning_data = pd.read_fwf('./ex7/drowning.txt').values
11 years = drowning_data[:, 0]
12 drowning = drowning_data[:, 1]
13
14
15 plt.plot(years, drowning)
16 z = np.polyfit(years, drowning, 1)
17 trend = np.poly1d(z)
18 plt.plot(years, trend(years), 'r--')
19
20 plt.savefig('./ex7/report/drowining.png')
21 plt.figure(0)
22
23 stan_code = '''
24 data {
25   int<lower=0> N; // number of data points
26   vector[N] x;   // observation year
```



```

27 vector[N] y;      // observation number of drowned
28 real xpred;      // prediction year
29 real tau;
30 }
31 parameters {
32   real alpha;
33   real beta;
34   real<lower=0> sigma;
35 }
36 transformed parameters {
37   vector[N] mu;
38   mu = alpha + beta * x;
39 }
40 model {
41   beta ~ normal(0, tau);
42   y ~ normal(mu, sigma);
43 }
44 generated quantities {
45   real ypred;
46   ypred = normal_rng(alpha + beta * xpred, sigma);
47 }
48 '''
49
50 ### guess of tau
51 dist = norm(loc=0, scale=20)
52 print(dist.cdf(-69))
53
54 ### fitting data to stan model
55 stan_model = pystan.StanModel(model_code=stan_code)
56
57 data = dict(
58     N=len(years),
59     x=years,
60     y=drowning,
61     xpred=2019,
62     tau=26.78,
63 )
64
65 ### sampling
66 fit = stan_model.sampling(data=data)
67 print(fit)
68
69 ### hist
70 y_pred = fit.extract()['ypred']
71 plt.hist(y_pred, bins=20, ec='white')
72 plt.show()

```

Appendix B Source code for Question 2

```

1 ###
2 import matplotlib

```

```

3  matplotlib.use('TkAgg')
4
5  from scipy.stats import norm
6  import matplotlib.pyplot as plt
7  import numpy as np
8  import pandas as pd
9  import pystan
10
11  ### The data
12  machines = pd.read_fwf('./ex7/factory.txt', header=None).values
13  machines_transposed = machines.T
14
15  ### Pooled model
16  '''
17  Pooled model
18  '''
19  stan_code_pooled = '''
20  data {
21      int<lower=0> N;          // number of data points
22      vector[N] y;           //
23  }
24  parameters {
25      real mu;                // group means
26      real<lower=0> sigma;    // common std
27  }
28  model {
29      y ~ normal(mu, sigma);
30  }
31  generated quantities {
32      real ypred;
33      ypred = normal_rng(mu, sigma);
34  }
35  '''
36
37  ### fitting data to stan model
38  machines_pooled = machines.flatten()
39  model_pooled = pystan.StanModel(model_code=stan_code_pooled)
40  data_pooled = dict(
41      N=machines_pooled.size,
42      y=machines_pooled
43  )
44
45  ### sampling
46  fit_pooled = model_pooled.sampling(data=data_pooled)
47  print(fit_pooled)
48
49  ### hist
50  y_pred_pooled = fit_pooled.extract()['ypred']
51  plt.hist(y_pred_pooled, bins=20, ec='white')
52  plt.savefig('./ex7/report/pooled_hist.png')
53  plt.figure(0)
54
55  mu = fit_pooled.extract()['mu']
56  plt.hist(mu, bins=20, ec='white')
57  plt.savefig('./ex7/report/pooled_hist_mu.png')

```

```

58 plt.figure(0)
59
60 ### Separate model
61 stan_code_separate = '''
62 data {
63     int<lower=0> N;                // number of data points
64     int<lower=0> K;                // number of groups
65     int<lower=1,upper=K> x[N];    // group indicator
66     vector[N] y;
67 }
68 parameters {
69     vector[K] mu;                 // group means
70     vector<lower=0>[K] sigma;     // group stds
71 }
72 model {
73     y ~ normal(mu[x], sigma[x]);
74 }
75 generated quantities {
76     real ypred;
77     ypred = normal_rng(mu[6], sigma[6]);
78 }
79 '''
80
81 ### fitting data into the stan model
82 model_separate = pystan.StanModel(model_code=stan_code_separate)
83 data_separate = dict(
84     N=machines_transposed.size,
85     K=6,
86     x=[
87         1, 1, 1, 1, 1,
88         2, 2, 2, 2, 2,
89         3, 3, 3, 3, 3,
90         4, 4, 4, 4, 4,
91         5, 5, 5, 5, 5,
92         6, 6, 6, 6, 6,
93     ],
94     y=machines_transposed.flatten()
95 )
96
97 ### sampling
98 fit_separate = model_separate.sampling(data=data_separate, n_jobs=-1)
99 print(fit_separate)
100
101 ### hist
102 y_pred_separate = fit_separate.extract()['ypred']
103 plt.hist(y_pred_separate, bins=20, ec='white')
104 plt.savefig('./ex7/report/separate_hist.png')
105 plt.figure(0)
106
107 ### hist
108 mu_data_separate = fit_separate.extract()['mu']
109 plt.hist(mu_data_separate[:, 5], bins=20, ec='white')
110 plt.savefig('./ex7/report/separate_hist_mu_six.png')
111 plt.figure(0)
112

```

```

113 ### Hierarchical model
114 stan_code_hierarchical = '''
115 data {
116     int<lower=0> N;           // number of data points
117     int<lower=0> K;           // number of groups
118     int<lower=1,upper=K> x[N]; // group indicator
119     vector[N] y;
120 }
121 parameters {
122     real mu0;                // prior mean
123     real<lower=0> sigma0;     // prior std
124     vector[K] mu;            // group means
125     real<lower=0> sigma;      // common std
126 }
127 model {
128     mu ~ normal(mu0, sigma0);
129     y ~ normal(mu[x], sigma);
130 }
131 generated quantities {
132     real ypred6;
133     real mu7;
134     ypred6 = normal_rng(mu[6], sigma);
135     mu7 = normal_rng(mu0, sigma0);
136 }
137 '''
138
139 ### fitting data into the stan model
140 model_hierarchical = pystan.StanModel(model_code=stan_code_hierarchical)
141 data_hierarchical = dict(
142     N=machines_transposed.size,
143     K=6,
144     x=[
145         1, 1, 1, 1, 1,
146         2, 2, 2, 2, 2,
147         3, 3, 3, 3, 3,
148         4, 4, 4, 4, 4,
149         5, 5, 5, 5, 5,
150         6, 6, 6, 6, 6,
151     ],
152     y=machines_transposed.flatten()
153 )
154
155 ### sampling
156 fit_hierarchical = model_hierarchical.sampling(data=data_hierarchical, n_jobs=-1)
157 print(fit_hierarchical)
158
159 ### hist
160 mu_data_hierarchical = fit_hierarchical.extract()['mu']
161 plt.hist(mu_data_hierarchical[:, 5], bins=20, ec='white')
162 plt.savefig('./ex7/report/hierarchical_hist_mu_six.png')
163 plt.figure(0)
164
165 ### hist
166 y_pred_hierarchical = fit_hierarchical.extract()['ypred6']
167 plt.hist(y_pred_hierarchical, bins=20, ec='white')

```

```
168 plt.savefig('./ex7/report/hierarchical_hist.png')
169 plt.figure(0)
170
171 ### hist
172 mu_data_hierarchical_7 = fit_hierarchical.extract()['mu7']
173 plt.hist(mu_data_hierarchical_7, bins=20, ec='white')
174 plt.savefig('./ex7/report/hierarchical_hist_mu_7.png')
175 plt.figure(0)
```